

Aishwarya Dudhe

Senior Data Engineer · Azure & Microsoft Fabric · Databricks & PySpark · Medallion Lakehouse · Streaming & Orchestration
+91-8483026896 | contact.adudhe@gmail.com | linkedin.com/in/aishwarya-dudhe | Bengaluru, India

PROFESSIONAL SUMMARY

Data platform engineer specialising in end-to-end Lakehouse ownership — Azure, Databricks, Snowflake, Microsoft Fabric — across financial services and enterprise manufacturing for Microsoft, NatWest Group, and HSBC. 7+ years designing distributed ETL/ELT pipelines, real-time streaming systems, and governed data platforms that serve 5,000+ end users at scale. Holds 5 active Microsoft and Databricks certifications.

CORE SKILLS

Languages: Python, SQL, PySpark, Spark SQL, Scala, T-SQL, KQL

Data Engineering: Distributed ETL/ELT, Delta Lake, Batch & Streaming (Spark, Kafka, Stream Analytics), Data Modeling (Star Schema, SCD2, Schema Evolution), Performance Tuning, Pipeline Scheduling & Monitoring, Data Quality Gates

Platforms: Databricks (Unity Catalog), Microsoft Fabric (Lakehouse, Pipelines, OneLake), Snowflake, Azure Synapse, Redshift, BigQuery, SQL Server

Cloud & Orchestration: Azure (ADF, ADLS Gen2, Stream Analytics, Fabric Pipelines, Key Vault), AWS (S3, Redshift), Microsoft Graph API, Git/CI-CD, Azure DevOps

BI & Governance: Power BI (Semantic Models, DAX, RLS, Paginated Reports), Tableau, RBAC, Data Governance, Data Lineage & Audit Logging

AI & RAG Infrastructure: RAG Architecture, Document Chunking & Embedding, Vector Databases (pgvector), Microsoft Graph API, Data Contracts

PROFESSIONAL EXPERIENCE

Senior Consultant – Data & AI | **Coretek** · Bengaluru, India

Jan 2026 – Present

Project: CAAP — Enterprise RAG-Enabled AI Agent Data Platform (Phase 1 MVP) · Jan 2026 – Apr 2026

- Architected the Phase 1 enterprise data backbone for a RAG-enabled AI agent platform on Microsoft Fabric — designing raw and curated Lakehouse zones for Gong, Salesforce, and SharePoint data with incremental ingestion, idempotency, and watermarking.
- Built SharePoint ingestion via Microsoft Graph API — crawling 15,800+ items across 9,200+ files and 6,600+ folders, with metadata-driven change detection (NEW/UPDATED/DELETED) and full version history tracking through a 5-stage registry-to-truth pipeline, orchestrated via Azure Data Factory on a daily incremental schedule.
- Owned the document processing pipeline design — defining chunking strategy, metadata schema, and pgvector indexing architecture for 9,200+ enterprise documents — enabling hybrid search retrieval and grounding for the Phase 1 AI agent.

Project: Laticrete International — Global Manufacturing & Construction Analytics · Apr 2026 – Present

- Architected end-to-end migration of 4 enterprise reports onto a unified Microsoft Fabric Medallion Lakehouse — consolidating 5 fragmented regional semantic models (INTL, MEA, SGP, NA) into a single governed star schema, cutting model maintenance overhead by ~60% and eliminating cross-region data inconsistencies.
- Engineered a Python-based .pbix transformation framework to programmatically rebind 1,000+ visuals across 80+ report pages — cutting a ~3-week manual migration cycle to hours and establishing a repeatable, version-controlled migration path deployed via Azure DevOps CI/CD.
- Designed a dynamic UPN-driven Row-Level Security architecture securing data access for ~250 users across 3 regional hierarchies — with fail-secure semantics ensuring zero unauthorized exposure on identity resolution failure, meeting enterprise audit requirements.
- Diagnosed a critical production pipeline drift (nightly job refreshing fact_sales without FX columns while the model was bound to fact_sales_revenue) — preventing silent revenue staleness post go-live and defining the P0 fix for nightly orchestration.

Data Engineer Consultant | **Version 1** · Bengaluru, India

May 2025 – Nov 2025

Client: NatWest Group — Financial Crime & Regulatory Reporting

- Architected and owned production Snowflake SQL pipelines ingesting financial crime datasets from upstream banking source systems — covering fraud detection, AML monitoring, and sanctions screening — feeding regulated reporting across 5+ compliance domains for one of the UK's largest banks.
- Automated Snowflake ETL orchestration with scheduled jobs, dependency chains, and failure alerting — reducing manual intervention by 40% and improving SLA adherence from ~85% to 98%+ on daily regulatory submissions consumed by compliance reporting dashboards.
- Tuned high-volume AML query workloads on Snowflake via clustering key redesign, warehouse right-sizing, and query profile analysis — cutting average scan runtime by ~50% (from ~60 min to ~30 min) and stabilizing nightly refresh windows.
- Enforced Git branch strategy and PR-based code reviews across all SQL pipeline changes — ensuring every regulatory transformation was version-controlled, peer-reviewed, and auditable, directly supporting compliance audit readiness.

Sr. Azure Data Engineer | [Lazydays RV](#) · Tampa, FL, USA

Oct 2023 – Dec 2024

\$1B+ revenue leisure vehicle retailer · 500+ employees nationwide · Full Azure platform ownership

- Owned the Azure data platform end-to-end — designed and operated ADF pipelines ingesting sales, inventory, accounting, and HR data from 8+ source systems into Azure SQL and Databricks, with data quality gates at every ingestion boundary sustaining a 99%+ data freshness SLA across all production feeds.
- Engineered the full ADF orchestration layer across 20+ production pipelines — triggers, dependency chains, retry logic, and failure alerting — and managed CI/CD deployments via Azure DevOps with Git-based branch strategy, serving as on-call engineer and reducing pipeline MTTR from ~4 hours to under 45 minutes.
- Built PySpark and Spark SQL transformations and star-schema dimensional models on Databricks with Delta Lake optimizations (partitioning, Z-ordering, OPTIMIZE/VACUUM) — delivering sub-15-min refresh cycles and reducing executive dashboard query latency by 35%.

Lead Data Engineer | [Microsoft \(via Infogain\)](#) · Seattle, WA, USA

Feb 2023 – Oct 2023

Contract Engagement — Enterprise HR Data Platform

- Designed ADF and KQL pipelines unifying HR data from 5+ source systems into a governed Azure Data Explorer model — enforcing data quality gates at ingestion, eliminating manual reconciliation, and cutting end-to-end data latency by 30% (~48 hrs to ~34 hrs).
- Implemented Azure Stream Analytics for real-time HR signal processing (joiners, attrition, asset events) — enabling workforce decisions ~2 days faster than the prior weekly batch approach.
- Delivered Power BI analytics adopted by 5,000+ HR employees for workforce planning — full rollout within 8 weeks with zero data security incidents, backed by RBAC, audit logging, and Key Vault secret management.

Data Engineer → Associate Consultant | [Capgemini](#) · Pune, India

Sep 2018 – Aug 2021

Clients: HSBC · HDFC Bank

Promoted 3 levels in 3 years: Software Engineer → Senior Software Engineer → Associate Consultant

- Built PySpark, Spark SQL, and Hive distributed pipelines on Databricks processing 2+ TB/day of transaction log data for HSBC's fraud analytics and AML risk stack — reducing batch processing time by 25% and supporting real-time compliance decisioning.
- Engineered ADF + Azure Synapse ETL across 10+ source systems into a centralized data warehouse — improving data availability 3x and reducing upstream-source onboarding time from ~3 weeks to under 1 week.
- Optimized PySpark transformations and partitioning strategies on Databricks — cutting processing time by 25% (from ~4 hours to ~3 hours on daily batch runs) and improving pipeline SLA adherence from ~88% to 96%+, directly strengthening AML regulatory compliance posture.
- Mentored 3 junior data engineers in PySpark best practices, pipeline design patterns, and Git-based development workflows — conducting weekly code reviews and defining team coding standards adopted across the Capgemini data practice.

CERTIFICATIONS

- Microsoft Certified: Azure Data Engineer Associate (DP-203)
- Microsoft Certified: Fabric Analytics Engineer Associate (DP-600)
- Microsoft Certified: Fabric Data Engineer Associate (DP-700)
- Databricks Certified Data Engineer Associate
- Databricks Certified Data Engineer Professional

EDUCATION

Master of Science, Business Analytics

2022

University of North Texas, Denton, TX, USA · GPA: 3.83 / 4.0 · Graduate Teaching Assistant, Data Visualization & BI (2021–2022)

Bachelor of Engineering, Electronics & Telecommunications

2018

Prof. Ram Meghe Institute of Technology & Research, Amravati, India · CGPA: 7.49 / 10.0